

# ReproBench: Benchmarking LLM Agents on CVE-only Vulnerability Reproduction

## Anonymous submission

### Abstract

Large language model (LLM) agents are increasingly evaluated on cybersecurity tasks, yet most benchmarks begin after the target environment has already been prepared: the agent receives source code, a patch, a running service, a container, or an executable binary. This setting leaves open a central question for real-world vulnerability analysis: can an agent reconstruct the missing environment needed to reproduce a known vulnerability from scratch? We introduce REPROBENCH, a public-evidence-grounded benchmark for from-scratch reproduction attempts of IoT firmware vulnerabilities from only a CVE identifier. REPROBENCH contains 30 manually curated IoT CVEs, balanced across buffer overflows, command injections, and authentication bypasses. Each task is accompanied by a manually inspected public-evidence dossier built from CVE records, NVD entries, vendor advisories, public PoCs, exploit databases, blogs, and firmware references. Agents must traverse a six-stage vulnerability reproduction pipeline: information gathering, firmware acquisition, firmware extraction, binary identification, service rehosting, and vulnerability triggering. We propose an evidence-grounded, non-binary scoring framework that separately evaluates planning and execution and requires observable artifacts for every non-zero score. Our evaluation consists of 450 reproduction runs (30 CVEs, 5 models, 3 trials) scored under a best-of-three protocol. Results show that agents average 50.1/100 overall, but only 4.7% of CVE-model pairs achieve real-target vulnerability triggering; 69.8% of runs substitute simulated reproductions for real firmware execution. REPROBENCH exposes the gap between textual vulnerability understanding and grounded firmware execution, revealing that the rehosting stage is the primary bottleneck where agents lose real-target grounding.

### Introduction

LLM agents now plan, call tools, write code, and operate in software repositories with increasing autonomy, producing benchmarks for software engineering and cybersecurity (Jimenez et al. 2024; Yang et al. 2026; Wang et al. 2025; Zhu et al. 2025; Lee et al. 2026; Zhang et al. 2026; Wang et al. 2026; Lee and Brumley 2026; Pu et al. 2026). A common assumption: the world has already been prepared for the

This is an anonymized submission for review purposes only. Distribution, citation, or public sharing of this manuscript is strictly prohibited. Copyright and publication details will appear in the final version if accepted.

agent. The agent receives a repository, patch, container, harness, or running service. Known-vulnerability reproduction in the wild begins in a less structured state: an analyst may start with only a CVE identifier and must discover the affected product, locate the correct artifact, recover a runnable environment, and validate vulnerability behavior against the real target. IoT firmware makes this gap especially sharp—firmware is a proprietary binary blob compiled for MIPS or ARM, packed in vendor-specific formats, dependent on missing NVRAM state, and reachable only after rehosting under emulation. Firmware reproduction is not merely exploit writing; it is an end-to-end artifact reconstruction and validation task.

We introduce REPROBENCH, a benchmark that evaluates whether LLM agents can move from a bare CVE identifier toward auditable evidence of IoT firmware vulnerability reproduction. REPROBENCH moves beyond prepared environments to evaluate the full artifact-to-execution pipeline, stressing capabilities that prepared-environment benchmarks hide: open-world information gathering, artifact provenance checking, cross-architecture binary analysis, long-horizon tool orchestration, and grounded validation. Each task is grounded in a manually inspected *public-evidence dossier*: a reference packet of public vulnerability facts and expected evidence types. Agents receive only the CVE identifier, not the dossier.

REPROBENCH also addresses a methodological problem: plausible-looking reproductions are easy to fabricate. When rehosting is difficult, agents produce substitutes such as Python mock servers or synthetic C programs that crash. REPROBENCH uses non-binary, evidence-grounded scoring across six stages with real-target gating for rehosting and triggering.

Our contributions are:

- We formulate CVE-only vulnerability reproduction as a from-scratch, long-horizon LLM agent benchmark.
- We construct a 30-CVE public-evidence dataset balanced across buffer overflow, command injection, and authentication bypass.
- We propose an evidence-grounded scoring framework with real-target gating across six dependent stages.
- We evaluate 450 runs: agents average 50.1/100 overall, but only 4.7% achieve real-target triggering; 69.8% involve

simulation substitution.

## Motivation and Task

Existing cybersecurity benchmarks measure post-environment capabilities: CyberGym and SEC-bench provide source-level projects (Wang et al. 2025; Lee et al. 2026); CVE-Bench provides prepared web targets (Zhu et al. 2025); ExploitBench evaluates after a target is available (Lee and Brumley 2026); K-Repro starts from kernel patch (Pu et al. 2026). None answer whether agents can create the reproduction environment themselves. REPROBENCH asks that question directly: the input is a CVE identifier only; the goal is auditable evidence of reproduction on the real affected firmware service. IoT firmware provides a compact stress test—the agent must retrieve public information, locate firmware that may have disappeared, distinguish real firmware from HTML errors, unpack vendor-specific formats, reason about non-x86 binaries, and use rehosting tools. Each stage leaves auditable artifacts.

REPROBENCH defines six dependent stages: P1 Information Gathering, P2 Firmware Acquisition, P3 Firmware Extraction, P4 Binary Identification, P5 Service Rehosting, and P6 Vulnerability Triggering. The dependency order is strict: P6 is valid only if the real service from P5 is reachable, and P5 is meaningful only after the correct firmware and binary have been identified. This structure is essential for rejecting simulation substitution.

## Benchmark Construction

### Manual CVE Curation

We construct REPROBENCH through manual curation from public CVE databases, vendor advisories, NVD entries, exploit databases, security blogs, and GitHub PoC repositories. A candidate is retained if it affects an IoT or embedded firmware product, public references identify the affected product or firmware family, the vulnerability class is assignable to our target categories, and there is enough public evidence to build a reference dossier. From the candidate pool, we select 30 CVEs balanced equally across three classes: 10 buffer overflows, 10 command injections, and 10 authentication bypasses. These classes require different validation signals—crashes for buffer overflows, command side effects for injection, and unauthenticated access for bypass—creating a controlled benchmark for comparing agent capabilities.

### Public-Evidence Dossiers

For each CVE, we construct a public-evidence dossier containing vendor and product information, affected firmware versions, vulnerability class, affected endpoint, public references, PoC descriptions, firmware acquisition hints, target binary names, architecture, and expected trigger evidence type. The dossier is used only for scoring; agents receive only the CVE identifier. The dossier is not a claim that the authors have independently reproduced every CVE end to end. Rather, it defines the publicly documented facts and evidence requirements against which an agent’s artifacts are audited.

## Evidence Scope and Attribution

The dossier is organized around the six stages. P1 fields describe checkable public facts; P2–P4 fields describe artifact expectations; P5–P6 fields specify acceptable evidence types for real execution (e.g., service banner, crash signal, command side effect) without providing an exploit oracle. This design separates three cases: agent capability gaps (fail despite strong evidence), evidence insufficiency (stale or missing public artifacts), and simulation substitution (plausible reports without dossier-supported artifacts). We use a conservative attribution rule: *artifact unavailable* requires confirmed stale sources plus a reasonable acquisition attempt; *agent acquisition failure* requires accessible sources but inadequate search; borderline cases are excluded from model-only failure counts.

## Evidence-Grounded Scoring

### Why Non-Binary Scoring?

Binary success is poorly suited to from-scratch firmware reproduction. Most runs fail before the final trigger, but a run that downloads and extracts the correct firmware has progressed much further than one that only reads an advisory. Conversely, an agent can fabricate a mock service and crash it, producing a convincing report without touching the real target. REPROBENCH therefore scores progress stage by stage and requires observable evidence for every non-zero score.

### Plan Score

Before execution, the agent writes a reproduction plan. The plan score has three components, each normalized to 0–100:  $\text{Plan} = 0.6 \times \text{Coverage} + 0.3 \times \text{Dependency} + 0.1 \times \text{Fallback}$ . *Coverage* evaluates whether the plan explicitly covers all six phases with the same weights as the task score. *Dependency correctness* checks that the planned order respects  $P1 \rightarrow P2 \rightarrow P3 \rightarrow P4 \rightarrow P5 \rightarrow P6$ , deducting 20 points per violation. *Fallback planning* is scored on a quality scale from 0 (no fallback) to 100 (phase-specific contingencies with blocked phase, failure mode, and alternative action).

### Task Score

The task score assigns 100 points across P1–P6: 15 points each for P1–P4 and 20 points each for P5–P6. Each phase is scored directly from its phase-specific evidence checklist; credit requires observable artifacts aligned with the public-evidence dossier. For example, P2 checks firmware acquisition, version verification (using version strings, meta-data, checksum or hash), and source address verification; P5 checks environment setup, binary launch, dependency initialization, service reachability, and real-firmware-service confirmation. Plan consistency, recovery attempts, and efficiency are recorded in failure attribution but do not receive separate task score points.

The overall score is  $\text{Overall} = 0.2 \cdot \text{Plan} + 0.8 \cdot \text{Task}$ . We weight execution more heavily because the benchmark measures whether an agent can realize a reproduction, not merely describe one.




Figure 1: Scoring framework overview. Each run is scored across six phases using evidence checklists, with real-target gating at P5/P6.

## Real-Target Gating

P5 and P6 enforce real-target validity: the reproduction task requires running the real vulnerable binary from the extracted firmware. Any approach that does not run the real firmware binary—mock servers, harnesses transcribing vulnerable logic, host-native simulations, or static-only analysis—is considered simulated reproduction and receives 0 for P5 and P6. This is not a penalty; the simulation work simply has no value for real-target reproduction. If the agent attempted real rehosting before switching to simulation, the real attempt is scored on its progress and the simulation is ignored.

## Scoring Skill and Audit

We use a rubric-constrained scoring skill to scale evaluation to 450 runs. The scorer receives the dossier, plan, trace, workspace, and reports for a single run, cites concrete evidence for every non-zero score, and outputs structured JSON. A stratified human audit validates the scoring on a sample of runs—focusing on non-zero P5/P6 claims, high scores, and simulation-like workspaces—but does not add hidden credit for unobserved work.

## Experimental Design

### Agent Runs

We evaluate five model configurations over the 30-CVE benchmark with three independent trials per CVE, for 450 reproduction runs: `claude-sonnet-4-6`, `deepseek-v4-flash-free`, `glm-5.2`, `gpt-5.5`, and `mimo-v2.5-free`. Each run receives a single CVE identifier, writes a plan, executes in an isolated workspace with a budget of \$10 and a 24-hour time limit, and records all artifacts. We preserve plans, command traces, downloaded files, extracted filesystems, scripts, logs, and reports. This

design deliberately allows public information use, including advisories and public PoCs; “from scratch” means the agent receives no prepared target environment, firmware image, or harness beyond the CVE identifier. A public PoC may help with payload syntax, but using it is not itself success—credit depends on grounding it in the correct firmware artifact and sending it to a real rehosted service.

## Scoring Runs

Every run is evaluated by the scoring skill, producing 450 JSON score files and summaries. For each CVE–model pair, we report the best plan score and the best task score across three trials (may come from different runs), combined as  $\text{Overall} = 0.2 \times \text{BestPlan} + 0.8 \times \text{BestTask}$ . This reflects the practical question: given multiple attempts, can an agent reproduce a vulnerability at least once? We also report failure attribution based on all 450 runs.

## Audit Trail

Each run is stored as an auditable record: prompt, plan, traces, downloaded files, extraction outputs, scripts, logs, reports, and scoring JSON. This design is important because invalid reproductions are only visible when the report is compared with the workspace—e.g., a report may claim a router endpoint crashed while the workspace shows a Python mock server. The scoring skill assigns stage-level credit; human auditors then validate evidence-citation consistency on a stratified sample.

We ask six research questions:

- RQ1: How far can agents progress from a CVE id toward reproduction?
- RQ2: Which stages account for most failures?
- RQ3: Does plan quality predict execution progress?
- RQ4: How often do agents produce plausible but invalid artifacts?

Table 1: Comparison with representative agent and cybersecurity benchmarks. Prepared environments hide the artifact-to-execution pipeline measured by REPROBENCH.

| Benchmark                           | Input          | Prepared env. | Artifact acquisition | Rehosting | Non-binary scoring   |
|-------------------------------------|----------------|---------------|----------------------|-----------|----------------------|
| SWE-bench (Jimenez et al. 2024)     | Repo + issue   | Yes           | No                   | No        | Tests                |
| ProgramBench (Yang et al. 2026)     | Binary + docs  | Partial       | No                   | No        | Test pass rate       |
| CyberGym (Wang et al. 2025)         | Source/harness | Yes           | No                   | No        | Staged               |
| CVE-Bench (Zhu et al. 2025)         | Web target     | Yes           | No                   | No        | Binary               |
| ExploitBench (Lee and Brumley 2026) | Built target   | Yes           | No                   | No        | Capability ladder    |
| K-Repro (Pu et al. 2026)            | Patch context  | Partial       | Limited              | Kernel    | Mostly final trigger |
| REPROBENCH                          | CVE ID         | No            | Yes                  | Yes       | Stage score          |

- RQ5: How do vulnerability classes and models affect progress?
- RQ6: Does evidence-grounded scoring reveal distinctions hidden by binary success?

## Evaluation Results

### Overall Capability (RQ1)

Table 2 reports the best-of-three results for each model across 30 CVEs. `glm-5.2` achieves the highest average overall score (65.2/100), followed by `mimo-v2.5-free` (50.3) and `claude-sonnet-4-6` (49.9). The best single run reaches 99.5/100 (`glm-5.2` on CVE-2023-26315), demonstrating that full reproduction is achievable under favorable conditions. However, strict real-target success remains rare: only 10 out of 150 CVE-model pairs achieve near-full P5 credit (score  $\geq 15/20$ ), and only 7 achieve near-full P6 credit.

The four CVEs that achieve full or near-full reproduction—CVE-2023-26315 (Xiaomi AX9000, 99.5), CVE-2023-44418 (95.0), CVE-2025-6443 (92.1), and CVE-2025-60690 (91.2)—share a common pattern: the target devices use ARM64 processors, allowing agents to run the real firmware binary via `chroot` without cross-architecture emulation. In contrast, MIPS-based devices requiring QEMU emulation almost universally fail at P5, because agents cannot reliably configure NVRAM, shared libraries, and network bridges.

A notable finding is that `glm-5.2`, not `claude-sonnet-4-6` or `gpt-5.5`, achieves the highest overall score. Per-run averages reveal the mechanism: `glm-5.2` averages 7.7/15 on P2 (Firmware Acquisition), nearly double `claude-sonnet-4-6`’s 4.5/15. Since P2 is a prerequisite for all downstream phases, this gap cascades—`claude-sonnet-4-6` has 17 runs where the plan scores above 50 but the task falls below 15, all blocked at P2 with zero downstream credit. Furthermore, `claude-sonnet-4-6` exhibits the highest simulation-substitution rate (81.1% of runs vs. `glm-5.2`’s 56.7%): when rehosting proves difficult, it rapidly pivots to writing mock servers or harnesses, producing polished reports that do not exercise the real firmware. `gpt-5.5` suffers from both weak firmware acquisition (P2 avg. 5.4) and 6 infrastructure failures (vs. 1 for `claude-sonnet-4-6`). In contrast, `glm-5.2` achieves the highest fallback planning score (37.2/100 vs. 25.8–31.7 for others), which correlates with its lower simulation rate and deeper pipeline

progression. These results suggest that on from-scratch reproduction tasks, the decisive factor is not raw model capability but task-appropriate behavior: persistent artifact acquisition, fallback-driven recovery, and resistance to simulation substitution.

**Result for RQ1:** Agents can progress partway from a CVE identifier toward vulnerability reproduction, averaging 50.1/100 under best-of-three scoring. However, real-target rehosting (P5) and triggering (P6) remain near-insurmountable barriers: only 6.7% of CVE-model pairs achieve near-full P5 credit and 4.7% achieve near-full P6. Full reproduction is achievable only on same-architecture devices where `chroot` suffices.

### Where Does the Pipeline Collapse? (RQ2)

Table 3 and Figure 2 show the stage funnel from P1 to P6. P1 (Information Gathering) is the strongest phase at 13.8/15, indicating that most agents successfully collect CVE facts. P2–P4 show moderate scores (7.0–8.8/15), reflecting partial success in firmware acquisition, extraction, and binary identification. The sharp collapse occurs at P5 (Service Rehosting), which averages only 2.7/20. P6 averages 2.6/20, with most points coming from PoC construction rather than real-target triggering.

Across all 450 runs, four failure families account for the terminal blockers. **Simulation substitution** dominates at 314/450 runs (69.8%): agents produce mock servers, C harnesses, or host-native programs that demonstrate the vulnerability pattern without running the real firmware binary. **Rehosting gap** accounts for 99/450 runs (22.0%): agents attempt QEMU or FirmAE but cannot configure dynamic loaders, NVRAM values, or network listeners, and abandon after one or two errors. **Artifact failures** (57/450, 12.7%) and **extraction/binary-analysis failures** (44/450, 9.8%) round out the taxonomy. Terminal blockers cluster at Phase 5 (246/450) and Phase 2 (143/450).

**Result for RQ2:** The pipeline collapses at P5 (Service Rehosting), which averages 2.7/20 even under best-of-three scoring. The dominant failure mode is simulation substitution (69.8% of all runs), followed by rehosting gap (22.0%). Terminal blockers concentrate at Phase 5 (54.7%) and Phase 2 (31.8%).

Table 2: Best-of-three model summary across 30 CVEs. Plan and Task are the best scores across three independent trials (may come from different runs). P5/P6 columns count CVEs with near-full credit ( $\geq 15/20$ ).

| Model                  | Avg Plan | Avg Task | Avg Overall | Max Overall | P5 $\geq$ 15 | P6 $\geq$ 15 |
|------------------------|----------|----------|-------------|-------------|--------------|--------------|
| glm-5.2                | 88.3     | 59.4     | <b>65.2</b> | 99.5        | 6/30         | 4/30         |
| mimo-v2.5-free         | 82.6     | 42.2     | 50.3        | 73.4        | 0/30         | 0/30         |
| claude-sonnet-4-6      | 80.4     | 42.2     | 49.9        | 99.0        | 2/30         | 1/30         |
| deepseek-v4-flash-free | 71.2     | 36.6     | 43.5        | 95.6        | 1/30         | 1/30         |
| gpt-5.5                | 71.2     | 34.3     | 41.7        | 91.8        | 1/30         | 1/30         |
| All                    | 78.8     | 43.0     | 50.1        | 99.5        | 10/150       | 7/150        |

Table 3: Stage-level averages under best-of-three scoring. The pipeline collapses at P5.

| Phase | Description              | Score | Max | %     |
|-------|--------------------------|-------|-----|-------|
| P1    | Information gathering    | 13.8  | 15  | 92.0% |
| P2    | Firmware acquisition     | 8.8   | 15  | 58.7% |
| P3    | Firmware extraction      | 8.0   | 15  | 53.3% |
| P4    | Binary identification    | 7.0   | 15  | 46.7% |
| P5    | Service rehosting        | 2.7   | 20  | 13.5% |
| P6    | Vulnerability triggering | 2.6   | 20  | 13.0% |

### Does Plan Quality Predict Execution? (RQ3)

We compare plan score against task score for each CVE-model pair. glm-5.2 achieves both the highest plan (88.3) and the highest task (59.4), with a gap of 28.9 points. gpt-5.5 has a plan of 71.2 but a task of only 34.3, a gap of 36.9. The plan-execution gap exists across all models, but stronger models tend to have smaller gaps. The fallback component is the weakest plan sub-score, averaging only 26.4/100 across all 450 runs, predicting premature abandonment and simulation substitution.

**Result for RQ3:** Plan quality partially predicts execution progress. glm-5.2 achieves both the highest plan (88.3) and task (59.4), with the smallest gap. However, the plan-execution gap persists across all models (29-37 points). Weak fallback planning (avg. 26.4/100) is the strongest predictor of simulation substitution.

### How Often Do Agents Simulate? (RQ4)

Simulation substitution is the most frequent failure mode: 314/450 runs (69.8%) produce simulated reproductions. Figure 3 shows per-phase progress rates across models. The rate varies significantly: claude-sonnet-4-6 substitutes in 81.1% of runs despite strong plan quality, while glm-5.2 substitutes in only 56.7%. Simulated outputs are often coherent—valid crash signals, correct exploit syntax, convincing reports—but never exercise the real firmware binary. Without real-target gating, these would be indistinguishable from successful reproduction.




Figure 2: The six-stage vulnerability reproduction pipeline. Agents receive a CVE number and must progress through each stage. The pipeline collapses at P5 (Service Rehosting), where cross-architecture emulation fails across nearly all runs.

**Result for RQ4:** 69.8% of runs involve simulation substitution. Stronger models do not necessarily substitute less—claude-sonnet-4-6 has the highest rate (81.1%)—suggesting that simulation is driven by task difficulty rather than model weakness. Real-target gating is essential to distinguish genuine reproduction from simulation.

### Vulnerability-Class and Model Analysis (RQ5)

Scores range from 99.5 (CVE-2023-26315) to 23.4 (CVE-2025-34037). The top-performing CVEs involve ARM64 devices where chroot suffices; the bottom involve MIPS devices with encrypted firmware or unavailable downloads. Seven CVEs show a >20-point gap between glm-5.2 and




Figure 3: Per-phase progress rates across models. P1–P4 show moderate success, but P5 and P6 collapse to near-zero for all models.

both `claude-sonnet-4-6` and `gpt-5.5`; in all seven, the weaker models fail at P2 (firmware acquisition) while `glm-5.2` downloads, extracts, and in some cases rehosts the firmware. Buffer overflows allow partial credit via static analysis (P4); command injections require real rehosting; authentication bypasses need stateful request comparisons. Vulnerability class affects the validation path but is secondary to rehosting feasibility.

**Result for RQ5:** CVE-level scores range from 99.5 to 23.4. Success correlates strongly with target architecture: ARM64 devices allowing `chroot`-based rehosting achieve full reproduction, while MIPS devices requiring QEMU almost universally fail at P5. Model-level differences are driven primarily by firmware acquisition success, not raw reasoning ability.

### Graded vs. Binary Metrics (RQ6)

Under binary evaluation, 143/150 pairs (95.3%) would be failures (P6 < 15/20). Graded scoring reveals a 23.5-point spread (41.7–65.2) and a stage funnel invisible to binary metrics. Best-of-three scoring adds ~10 points over per-run averages, with stronger models showing larger variance (e.g., `glm-5.2`: +14.3), suggesting that succeeding at least once is itself a capability dimension.

**Result for RQ6:** Evidence-grounded scoring reveals distinctions hidden by binary success. Binary evaluation classifies 95.3% as failures, but graded scoring shows a 23.5-point spread. The stage funnel, simulation-substitution rate, and plan-execution gap are all invisible under binary metrics.

### Failure Modes

Section quantifies four failure families across 450 runs, illustrated by the case study in Figure 4.

**Artifact failures (57/450, 12.7%).** Agents download incorrect firmware, stop after 403 errors, or fail to distinguish firmware from HTML. Terminal blockers at Phase 2 account for 143/450 runs.

**Extraction and binary-analysis failures (44/450, 9.8%).** Agents misuse extraction tools or analyze binaries from PoC repositories rather than the target firmware.

### Case Study: CVE-2020-13389 (Tenda AC15, MIPS32-LE, Stack Overflow)




Figure 4: Case study: the agent succeeds through P4 but substitutes a simulated reproduction at P5/P6.

**The rehosting gap (99/450, 22.0%).** Agents know QEMU and firmware emulators but cannot configure dynamic loaders, NVRAM, or network listeners, and abandon after one or two errors. Same-architecture (ARM64) devices bypass this gap via `chroot`.

**Simulation substitution (314/450, 69.8%).** The dominant failure: agents produce mock servers, C harnesses, or host-native programs that demonstrate the vulnerability pattern but not the instance. The output may be coherent and reproducible, but wrong for the benchmark goal. Real-target gating rejects all of these.

## Discussion

### What REPROBENCH Measures

REPROBENCH evaluates whether agents can produce auditable evidence aligned with public vulnerability facts and real firmware artifacts, without assuming a prebuilt executable oracle for each CVE.

### Implications for Agent Evaluation

Prepared environments hide a major class of agent failures. Future benchmarks should record traces and workspaces, require artifact provenance, distinguish evidence insufficiency from agent error, and reject substitutions that do not preserve task validity.

### Implications for Agent Design

The results reveal that the decisive factor in from-scratch reproduction is not raw model capability but task-appropriate behavior. `glm-5.2` outperforms `claude-sonnet-4-6` and `gpt-5.5` despite being a smaller model, because it persists longer at firmware acquisition (P2 avg. 7.7 vs. 4.5),



plans better fallbacks (37.2 vs. 25.8–31.7), and resists simulation substitution (56.7% vs. 81.1%). This suggests that future agent designs should prioritize persistent artifact acquisition, fallback-driven recovery, and self-checks that distinguish real targets from simulations. The architecture-dependent ceiling—same-architecture (ARM64) devices allow `chroot`-based rehosting while cross-architecture (MIPS, ARM32) devices almost universally fail—indicates that cross-architecture emulation configuration (NVRAM initialization, library path resolution, network bridge setup) is the specific blocker agents must overcome.

## Limitations and Ethics

The dataset is balanced across three vulnerability classes but does not mirror the natural IoT distribution. The scoring skill is not a deterministic exploit oracle; human analysis validates evidence consistency. Ethically, REPROBENCH targets defensive reproduction using only publicly disclosed, remediated CVEs, conducted in isolated workspaces.

## Related Work

Software and agent benchmarks such as SWE-bench and ProgramBench evaluate code repair or program reconstruction (Jimenez et al. 2024; Yang et al. 2026). REPROBENCH shares ProgramBench’s from-scratch philosophy but focuses on reconstructing vulnerability reproduction environments.

Cybersecurity benchmarks evaluate web exploitation, CTF, vulnerability reproduction, patching, and exploit construction (Wang et al. 2025; Zhu et al. 2025; Lee et al. 2026; Zhang et al. 2026; Wang et al. 2026; Lee and Brumley 2026; Pu et al. 2026). Most provide source code, patch context, a harness, or a prepared service. REPROBENCH complements them by evaluating the artifact-to-execution pipeline that precedes exploitation in IoT firmware reproduction.

Firmware rehosting tools such as Firmadyne, FirmAE, and Greenhouse automate parts of emulation (Chen et al. 2016; Kim et al. 2020; Tay et al. 2023). REPROBENCH does not replace these tools; it evaluates whether LLM agents can discover, install, configure, and recover from failures while using them in a from-scratch workflow.

## Conclusion

REPROBENCH evaluates whether LLM agents can move from a CVE identifier to auditable evidence of IoT firmware vulnerability reproduction. Across 450 runs under best-of-three scoring, agents average 50.1/100, with strong performance through information gathering and firmware analysis, but collapse at rehosting (2.7/20). 69.8% of runs substitute simulated reproductions; real-target success is achievable only on same-architecture devices where `chroot` bypasses emulation. Future cyber agents must be evaluated not only on plausible reports, but on whether their actions remain grounded in the real artifacts they claim to reproduce.

## References

Chen, D. D.; Woo, M.; Brumley, D.; and Egele, M. 2016. Towards automated dynamic analysis for linux-based embedded firmware. In *NDSS*, volume 1, 1–1.

Jimenez, C. E.; Yang, J.; Wettig, A.; Yao, S.; Pei, K.; Press, O.; and Narasimhan, K. 2024. Swe-bench: Can language models resolve real-world github issues? In *International Conference on Learning Representations*, volume 2024, 54107–54157.

Kim, M.; Kim, D.; Kim, E.; Kim, S.; Jang, Y.; and Kim, Y. 2020. Firmae: Towards large-scale emulation of iot firmware for dynamic analysis. In *Proceedings of the 36th Annual Computer Security Applications Conference*, 733–745.

Lee, H.; Zhang, Z.; Lu, H.; and Zhang, L. 2026. Sec-bench: Automated benchmarking of llm agents on real-world software security tasks. *Advances in Neural Information Processing Systems*, 38: 116342–116378.

Lee, S.; and Brumley, D. 2026. ExploitBench: A Capability Ladder Benchmark for LLM Cybersecurity Agents. *arXiv preprint arXiv:2605.14153*.

Pu, J.; Li, X.; Liang, Z.; Cox, J.; Wu, Y.; Shehada, K.; Srivastav, A.; and Qian, Z. 2026. Patch-to-PoC: A Systematic Study of Agentic LLM Systems for Linux Kernel N-Day Reproduction. *arXiv preprint arXiv:2602.07287*.

Tay, H. J.; Zeng, K.; Vadayath, J. M.; Raj, A. S.; Dutcher, A.; Reddy, T.; Gibbs, W.; Basque, Z. L.; Dong, F.; Smith, Z.; et al. 2023. Greenhouse: {single-service} rehosting of {linux-based} firmware binaries in {user-space} emulation. In *32nd USENIX Security Symposium (USENIX Security 23)*, 5791–5808.

Wang, Z.; Schiller, N.; Li, H.; Narayana, S. S.; Nasr, M.; Carlini, N.; Qi, X.; Wallace, E.; Bursztein, E.; Invernizzi, L.; et al. 2026. ExploitGym: Can AI Agents Turn Security Vulnerabilities into Real Attacks? *arXiv preprint arXiv:2605.11086*.

Wang, Z.; Shi, T.; He, J.; Cai, M.; Zhang, J.; and Song, D. 2025. CyberGym: Evaluating AI Agents’ Real-World Cybersecurity Capabilities at Scale. *arXiv preprint arXiv:2506.02548*.

Yang, J.; Lieret, K.; Ma, J.; Thakkar, P.; Pedchenko, D.; Sootla, S.; McMilin, E.; Yin, P.; Hou, R.; Synnaeve, G.; et al. 2026. ProgramBench: Can Language Models Rebuild Programs From Scratch? *arXiv preprint arXiv:2605.03546*.

Zhang, A.; Ji, J.; Menders, C.; Dulepet, R.; Qin, T.; Wang, R.; Wu, J.; Liao, K.; Li, J.; Hu, J.; et al. 2026. Bounty-bench: Dollar impact of ai agent attackers and defenders on real-world cybersecurity systems. *Advances in Neural Information Processing Systems*, 38.

Zhu, Y.; Kellermann, A.; Bowman, D.; Li, P.; Gupta, A.; Danda, A.; Fang, R.; Jensen, C.; Ihli, E.; Benn, J.; et al. 2025. CVE-bench: a benchmark for AI agents’ ability to exploit real-world web application vulnerabilities. *arXiv preprint arXiv:2503.17332*.